

```
In [8]: import tensorflow as tf # tensorflow 2.0
from keras.datasets import mnist
import numpy as np
seed=0
np.random.seed(seed) # fix random seed
tf.random.set_seed(seed)
# input image dimensions
num_classes = 10 # 10 digits

img_rows, img_cols = 28, 28 # number of pixels

# the data, shuffled and split between train and test sets
(X_train, Y_train), (X_test, Y_test) = mnist.load_data()
```

```
In [9]: X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)
# cast floats to single precision
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
# rescale data in interval [0,1]
X_train /= 255
X_test /= 255
```

```
In [14]: Y_train = keras.utils.np_utils.to_categorical(Y_train, num_classes)
Y_test = keras.utils.np_utils.to_categorical(Y_test, num_classes)
```

```
In [15]: from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten
from keras.layers import MaxPooling2D, Dropout
model = Sequential()#add model layers
model.add(Conv2D(32, kernel_size=(5, 5),
                activation='relu',
                input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2)))
# add second convolutional layer with 20 filters
model.add(Conv2D(64, (5, 5), activation='relu'))

# add 2D pooling layer
model.add(MaxPooling2D(pool_size=(2, 2)))

# flatten data
model.add(Flatten())

# add a dense all-to-all relu layer
model.add(Dense(1024, activation='relu'))

# apply dropout with rate 0.5
model.add(Dropout(0.5))

# soft-max layer
model.add(Dense(num_classes, activation='softmax'))
```

```
In [16]: #compile model using accuracy to measure model performance
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [17]: #train the model
model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=3)
```

```
Epoch 1/3
1875/1875 [=====] - 95s 50ms/step - loss: 0.1196 - accuracy: 0.9635 - val_loss: 0.0320 - val_accuracy: 0.9891
Epoch 2/3
1875/1875 [=====] - 93s 50ms/step - loss: 0.0430 - accuracy: 0.9863 - val_loss: 0.0240 - val_accuracy: 0.9919
Epoch 3/3
1875/1875 [=====] - 93s 50ms/step - loss: 0.0311 - accuracy: 0.9901 - val_loss: 0.0332 - val_accuracy: 0.9901
```

Out[17]: <keras.callbacks.History at 0x7f76dc7d1850>

```
In [18]: # evaluate the model
score = model.evaluate(X_test, Y_test, verbose=1)

# print performance
print()
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
313/313 [=====] - 5s 16ms/step - loss: 0.0332 - accuracy: 0.9901
```

```
Test loss: 0.03317980840802193
Test accuracy: 0.9901000261306763
```

```
In [21]: predict_x=model.predict(X_test[:4])
classes_x=np.argmax(predict_x,axis=1)
```

```
In [22]: #predict first 4 images in the test set
model.predict(X_test[:4])
```

```
Out[22]: array([[4.44570977e-11, 7.88449361e-10, 1.57920912e-08, 1.73734200e-08,
 1.06866394e-09, 4.25641689e-10, 1.14703603e-13, 9.99999881e-01,
 2.63262634e-10, 1.12674442e-07],
 [6.04240370e-07, 1.02885814e-07, 9.99999285e-01, 6.81972256e-10,
 2.71400263e-10, 1.80242770e-13, 1.30808320e-09, 1.46937351e-10,
 1.58902793e-08, 1.04200828e-12],
 [2.68074842e-08, 9.99998927e-01, 1.80928353e-07, 1.01474285e-09,
 5.21466355e-08, 8.92409702e-09, 4.40140866e-08, 5.91038884e-07,
 1.31291557e-07, 3.73260489e-09],
 [9.99999166e-01, 5.14480236e-10, 6.72057467e-07, 1.56106472e-09,
 3.76253917e-10, 4.15453255e-10, 2.55185695e-09, 5.21882138e-09,
 1.13521956e-08, 9.95265168e-08]], dtype=float32)
```

```
In [24]: #actual results for first 4 images in test set
Y_test[:4]
```

```
Out[24]: array([[0., 0., 0., 0., 0., 0., 0., 1., 0., 0.],
 [0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```

In [ ]: