

```
In []: import pandas as pd
pd.read_csv('../input/traffic-signs-classification/labels.csv')
```

Importing Libraries

```
In [2]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.image import imread
import seaborn as sns
```

Data Collection

```
In [3]: dir_path = '../input/gtsrb-german-traffic-sign'
```

```
In [4]: os.listdir(dir_path)
```

```
Out[4]: ['Meta',
'meta',
'Meta.csv',
'Train.csv',
'Test.csv',
'Test',
'test',
'Train',
'train']
```

```
In [5]: #Assigning the path for train and test images
```

```
train_path=dir_path+'Train'
test_path=dir_path+'Test'
```

```
In [6]: print(sorted(os.listdir(train_path)))
```



```
[0', '1', '10', '11', '12', '13', '14', '15', '16', '17', '18', '19', '2', '20', '21', '22',  
'23', '24', '25', '26', '27', '28', '29', '3', '30', '31', '32', '33', '34', '35', '36', '3  
7', '38', '39', '4', '40', '41', '42', '5', '6', '7', '8', '9']
```

```
In []: sorted(os.listdir(test_path))
```

Visualization

Visualizing 25 random sample images from test set

```
In [8]: import random  
images_path = os.listdir(test_path)  
plt.figure(figsize=(25,25))  
  
for i in range(1,26):  
  
    plt.subplot(5,5,i)  
    random_img_path = test_path + '/' + random.choice(images_path) random_img =  
    imread(random_img_path)  
    plt.imshow(rand_img)  
    plt.xlabel(rand_img.shape[1], fontsize = 20) #width of image  
    plt.ylabel(rand_img.shape[0], fontsize = 20) #height of image
```





The dimensions of the images are not fixed.

Note:

Convolutional neural networks cannot perform on images that have various dimensions. We will resize these images during our model building.


But first find the mean of the dimensions of all the images in training set.

```
In [9]: dim1=[]
dim2=[]

for i in range(0,43):
    labels=train_path+'/{0}'.format(i)
    image_paths=os.listdir(labels)
    for image_path in image_paths:
        img=imread(labels+'/'+image_path)
        dim1.append(img.shape[0])
        dim2.append(img.shape[1])
```

Exploring the dimensions with a jointplot

```
In [10]: sns.jointplot(dim1,dim2).plt.show()
```



In [11]: `np.mean(dim1)`

Out[11]: 50.328929582493814

In [12]: `np.mean(dim2)`

Out[12]: 50.83587951745773

Since the mean of both dimensions is around 50 , we will use (50x50) as the shape of images.

In [13]: `image_shape=(50,50)`

Data Preprocessing

Importing the images

```
In [14]: from PIL import Image

images=[]
label_id=[]

for i in range(43):
    labels=train_path+'/{0}'.format(i)
    image_path=os.listdir(labels)
    for x in image_path:
        img=Image.open(labels+'/'+x)
        img=img.resize(image_shape)
        img=np.array(img)
        images.append(img)
```

Scaling the images so that the values remain between 0 and 1

```
In [15]: #Converting images into numpy array
images=np.array(images)
#The pixel value of each image ranges between 0 and 255
#Dividing each image by 255 will scale the values between 0 and 1. This is also known as norm
images=images/255
```


```
In [16]: label_id=np.array(label_id)
label_id.shape
```

Out[16]: (39209,)

```
In [17]: images.shape
```

Out[17]: (39209, 50, 50, 3)

```
In [18]: plt.figure(figsize=(15,5))
sns.countplot(label_id)
plt.title('Distribution of images among different classes',fontsize=15)
plt.xlabel('Label_id',fontsize=12)
plt.show()
```



```
In [19]: #Saving the scaled images and labels for future use
np.save('Training_set',images)
np.save('Label_Id',label_id)
```

Splitting the train data into train and validation data

```
In [20]: import numpy as np
import pandas as pd
```

```
In [21]: images=np.load('Training_set.npy')
label_id=np.load('Label_Id.npy')
```

```
In [22]: #Splitting the data
from sklearn.model_selection import train_test_split
x_train,x_val,y_train,y_val=train_test_split(images,label_id,test_size=0.2,random_
```

Changing target labels to categorical using one-hot encoding technique

```
In [23]: #keras has a built-in function for one-hot encoding.
from tensorflow.keras.utils import to_categorical

y_train = to_categorical(y_train)
y_val = to_categorical(y_val)
```



Model Building

```
In [24]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout, Conv2D, MaxPool2D
```

```
In [25]: model = Sequential()

#1st layer
model.add(Conv2D(filters=64, kernel_size=(3,3), input_shape=x_train.shape[1:], activation='relu'))
model.add(MaxPool2D(pool_size=(2,2)))
model.add(Dropout(0.5))

#2nd layer
model.add(Conv2D(filters=64, kernel_size=(3,3), activation='relu'))
model.add(MaxPool2D(pool_size=(2,2)))
model.add(Dropout(0.5))

#3rd layer
model.add(Conv2D(filters=64, kernel_size=(3,3), activation='relu'))
model.add(MaxPool2D(pool_size=(2,2)))
model.add(Dropout(0.5))
model.add(Flatten())

#Dense layer
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))

#Output layer
model.add(Dense(43, activation='softmax'))
```

```
In [26]: model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [27]: model.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 50, 50, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 25, 25, 64)	0
dropout (Dropout)	(None, 25, 25, 64)	0
conv2d_1 (Conv2D)	(None, 23, 23, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 64)	0
dropout_1 (Dropout)	(None, 11, 11, 64)	0
conv2d_2 (Conv2D)	(None, 9, 9, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_2 (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 43)	5547
=====		
Total params: 212,395		
Trainable params: 212,395		
Non-trainable params: 0		

In [28]:

```
early_stopping=tf.keras.callbacks.EarlyStopping(monitor='val_loss',patience=2)model.fit(

    x_train,y_train,epochs=25
    ,
    batch_size=64,
    validation_data=(x_val,y_val),callbac
    ks=[early_stopping],
    verbose=2
)
```



```

Epoch 1/25
491/491 - 85s - loss: 2.6929 - accuracy: 0.2470 - val_loss: 1.7082 - val_accuracy: 0.5179 Epoch 2/25
491/491 - 83s - loss: 1.4750 - accuracy: 0.5205 - val_loss: 0.8227 - val_accuracy: 0.7754 Epoch 3/25
491/491 - 83s - loss: 0.9164 - accuracy: 0.6967 - val_loss: 0.3744 - val_accuracy: 0.9218 Epoch 4/25
491/491 - 83s - loss: 0.6258 - accuracy: 0.7922 - val_loss: 0.2281 - val_accuracy: 0.9536 Epoch 5/25
491/491 - 83s - loss: 0.4927 - accuracy: 0.8393 - val_loss: 0.1606 - val_accuracy: 0.9688 Epoch 6/25
491/491 - 83s - loss: 0.4136 - accuracy: 0.8658 - val_loss: 0.1222 - val_accuracy: 0.9733 Epoch 7/25
491/491 - 82s - loss: 0.3639 - accuracy: 0.8815 - val_loss: 0.1012 - val_accuracy: 0.9821 Epoch 8/25
491/491 - 85s - loss: 0.3263 - accuracy: 0.8935 - val_loss: 0.0797 - val_accuracy: 0.9824 Epoch 9/25
491/491 - 84s - loss: 0.2886 - accuracy: 0.9071 - val_loss: 0.0650 - val_accuracy: 0.9866
Epoch 10/25
491/491 - 84s - loss: 0.2714 - accuracy: 0.9128 - val_loss: 0.0576 - val_accuracy: 0.9876
Epoch 11/25
491/491 - 83s - loss: 0.2599 - accuracy: 0.9159 - val_loss: 0.0547 - val_accuracy: 0.9890
Epoch 12/25
491/491 - 84s - loss: 0.2444 - accuracy: 0.9226 - val_loss: 0.0559 - val_accuracy: 0.9901
Epoch 13/25
491/491 - 84s - loss: 0.2290 - accuracy: 0.9260 - val_loss: 0.0456 - val_accuracy: 0.9888
Epoch 14/25
491/491 - 83s - loss: 0.2220 - accuracy: 0.9292 - val_loss: 0.0399 - val_accuracy: 0.9923
Epoch 15/25
491/491 - 84s - loss: 0.2234 - accuracy: 0.9292 - val_loss: 0.0472 - val_accuracy: 0.9895
Epoch 16/25
491/491 - 87s - loss: 0.2015 - accuracy: 0.9369 - val_loss: 0.0342 - val_accuracy: 0.9921
Epoch 17/25
491/491 - 84s - loss: 0.1986 - accuracy: 0.9371 - val_loss: 0.0354 - val_accuracy: 0.9922
Epoch 18/25
491/491 - 84s - loss: 0.1904 - accuracy: 0.9396 - val_loss: 0.0329 - val_accuracy: 0.9934
Epoch 19/25
491/491 - 83s - loss: 0.1858 - accuracy: 0.9411 - val_loss: 0.0370 - val_accuracy: 0.9921
Epoch 20/25
491/491 - 84s - loss: 0.1785 - accuracy: 0.9429 - val_loss: 0.0300 - val_accuracy: 0.9946
Epoch 21/25
491/491 - 84s - loss: 0.1808 - accuracy: 0.9439 - val_loss: 0.0267 - val_accuracy: 0.9950
Epoch 22/25
491/491 - 83s - loss: 0.1784 - accuracy: 0.9447 - val_loss: 0.0272 - val_accuracy: 0.9946
Epoch 23/25
491/491 - 87s - loss: 0.1657 - accuracy: 0.9487 - val_loss: 0.0280 - val_accuracy: 0.9934

```

Out[28]: <tensorflow.python.keras.callbacks.History at 0x7f2a98132fd0>

Achieved highest accuracy of 99.50% on validation data

```

In [29]: #Saving the model
model.save('Model.h5')

```

Model Evaluation

```

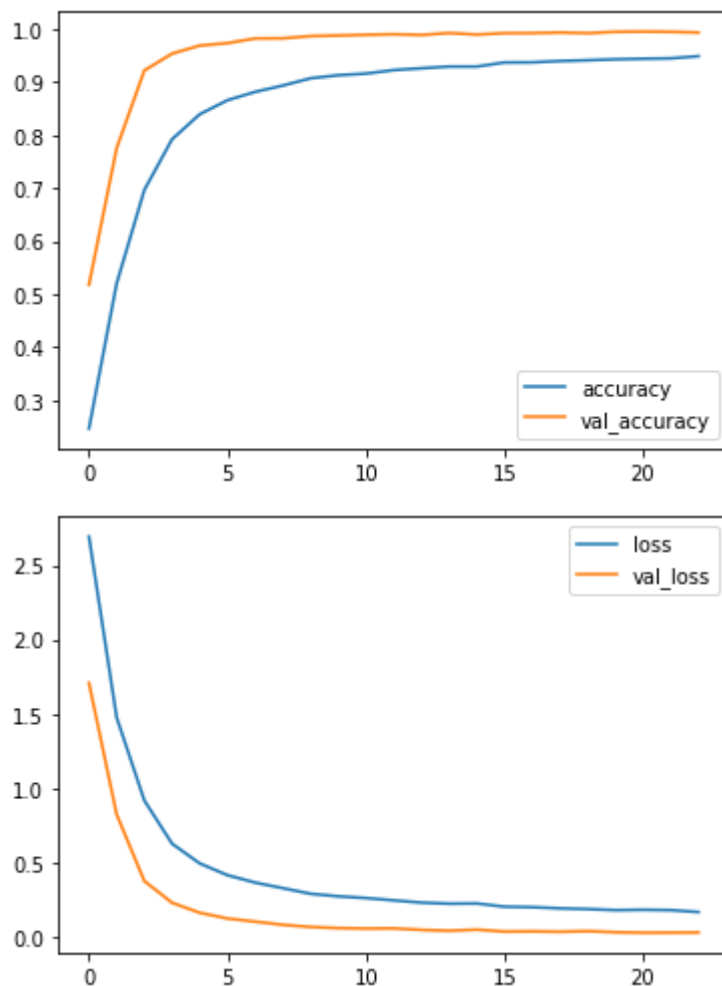
In [30]: evaluation=pd.DataFrame(model.history.history)

evaluation[['accuracy','val_accuracy']].plot()evaluation[['loss','val_loss']].plot()

```

Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2a90091610>





Testing on test data

In [31]: `from tensorflow.keras.models import load_model
model = load_model('Model.h5')`

Note: The test images folder in the original dataset has a blank csv file which cannot be opened with the above function. So i copied that folder and deleted that csv file and uploaded the test images again separately. These test images are same as the test images in the original dataset

In [32]: `test_path = '../input/test-images/Test'
test_img = sorted(os.listdir(test_path))`

For some unknown reason , the images in kaggle kernel is not showing in the order they are in the Test folder. Upon inspection it is seen that the images are in sorted order. So using sorted() function to sort them.

In [33]: `#defining a function that will scale images
from PIL import Image

def scaling(test_images, test_path):
 images = []

 for image_path in test_images:
 img = Image.open(test_path + '/' + image_path)
 img = img.resize((50, 50))
 img = np.array(img)
 images.append(img)`



```
#Converting images into numpy array
images=np.array(images)
#The pixel value of each image ranges between 0 and 255
#Dividing each image by 255 will scale the values between 0 and 1. This is also known as
images=images/255

return images
```

The above function can be used to scale any new traffic-sign images that can be predicted with our model. This is a general purpose function for code reusability.

```
In [34]: test_images=scaling(test_img,test_path)
```

Test labels

```
In [35]: test=pd.read_csv('../input/gtsrb-german-traffic-
sign/Test.csv')y_test=test['ClassId'].values
```

```
Out[35]: y_test
```

Testing on test images

```
In [36]: y_pred=model.predict_classes(test_images)

y_pred
```

```
Out[36]: array([16, 1, 38, ..., 6, 7, 10])
```

```
In [37]: from sklearn.metrics import classification_report

print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	1.00	0.98	0.99	60
1	0.99	0.98	0.99	720
2	0.99	0.99	0.99	750
3	0.99	0.93	0.96	450
4	0.98	0.99	0.99	660
5	0.94	0.97	0.96	630
6	1.00	0.93	0.96	150
7	0.98	0.98	0.98	450
8	0.97	0.98	0.98	450
9	0.95	1.00	0.97	480
10	0.99	1.00	1.00	660
11	0.97	1.00	0.98	420
12	0.96	0.94	0.95	690
13	0.99	1.00	1.00	720
14	1.00	1.00	1.00	270
15	0.90	1.00	0.95	210
16	1.00	1.00	1.00	150
17	1.00	0.92	0.96	360
18	0.99	0.95	0.97	390
19	0.92	1.00	0.96	60
20	0.72	0.99	0.83	90
21	0.94	0.64	0.76	90
22	0.96	0.97	0.97	120
23	0.99	0.96	0.98	150
24	1.00	0.97	0.98	90
25	0.96	0.96	0.96	480
26	0.95	0.90	0.93	180
27	0.80	0.47	0.59	60
28	0.97	1.00	0.99	150
29	0.74	0.99	0.85	90
30	1.00	0.77	0.87	150
31	0.96	0.94	0.95	270
32	0.85	1.00	0.92	60
33	0.97	0.99	0.98	210
34	0.98	0.99	0.99	120
35	1.00	0.99	1.00	390
36	0.96	0.99	0.98	120
37	0.98	0.98	0.98	60
38	0.95	0.99	0.97	690
39	0.98	0.99	0.98	90
40	0.98	0.99	0.98	90
41	1.00	0.58	0.74	60
42	0.96	1.00	0.98	90
accuracy			0.97	12630
macro avg	0.96	0.94	0.95	12630
weighted avg	0.97	0.97	0.97	12630

We achieved an overall accuracy of 97% on our model. This is pretty good and we can use this model for predicting some other Traffic signs as well in future.

